


Car Detection Project

In this report I will be sharing the details of an object detection model created using azure custom vision service. The name of the model is car detector which uses the images of different cars on a highway to detect different cars and trucks.

Model Creation

In order to create an efficient model, I used the custom vision AI service and uploaded the required data of 50 images. I tagged these images so that the model can understand the different types of objects present in the images. Once the tagging was complete I trained the model and achieved a respectable score. Below are the screenshots of the data, tags and model performance.



Tags +

Tagged Untagged

Showing: all tagged images

Search For Tags:

car 54 ...

truck 32 ...

Training Images Performance Predictions Train Quick Test ?

Unpublish Prediction URL Delete Export

s223124219@deakin.edu.au

Finished training on 4/22/2024, 5:43:42 PM using General domain
Iteration id: 35eb606d-f9a1-4d10-b3f3-34e5f549f650
Published as: CarDetector

Precision 81.4% Recall 61.4% mAP 55.8%

Utilizing the model using python code

Once the model training was completed, I imported all the required libraries and passed in the required credentials to gain access to my azure account.

```

1 from azure.cognitiveservices.vision.customvision.prediction import CustomVisionPredictionClient
2 from msrest.authentication import ApiKeyCredentials
3 import matplotlib.pyplot as plt
4 from PIL import Image, ImageDraw, ImageFont
5 import numpy as np
6 import os
7 from sort import Sort
8 import cv2

```

In order to gain access to azure account, I had to pass in the prediction endpoint, prediction key, project id and model name. This allows us to use the model in our code that we have trained using cognitive custom vision services.

Custom Vision

Settings

Resources:

Create New

cvtskd-Prediction

Subscription: Azure for Students
Resource Group: tsk_5d
Resource Kind: Custom Vision Prediction

Key:
6e77197d794e43f28b413422f01b7d47

Endpoint:
https://cvtskd-prediction.cognitiveservices.azure.com/

Resource Id:
/subscriptions/8137f162-de83-4bc5-adbc-7a3c0d0e4b59/resourceGroups/tsk_5d/providers/Microsoft.CognitiveServices/

Sign out

CarDetector

Training Images Performance Predictions Train Quick Test ?

Project Settings

General

Project Name*
CarDetector

Project Id
2368dec4-29b3-4cfc-a78b-ce469bba6f5d

Settings

Resources:

cvtskd

Subscription: Azure for Students
Resource Group: tsk_5d
Resource Kind: Custom Vision Training

Sign out

```

1 prediction_endpoint = 'https://cvtskd-prediction.cognitiveservices.azure.com/'
2 prediction_key = '6e77197d794e43f28b413422f01b7d47'
3 project_id = '2368dec4-29b3-4cfc-a78b-ce469bba6f5d'
4 model_name = 'CarDetector'

1 # Authenticate a client for the training API
2 credentials = ApiKeyCredentials(in_headers={"Prediction-key": prediction_key})
3 prediction_client = CustomVisionPredictionClient(endpoint=prediction_endpoint, credentials=credentials)

```

Credentials variable is used to create authentication credentials for accessing the Custom Vision prediction service. It utilizes the ApiKeyCredentials class from the azure.core.credentials module and passes the prediction key as a header. Prediction_client endpoint variable stores the endpoint URL for accessing the Custom Vision prediction service. The endpoint URL points to the Custom Vision service provided by Azure Cognitive Services.

Test

Before applying the object tracking, I decided to first test the model with one of the images to see if it is accurately detecting different objects in one of the images.

```

In [62]: 1 image_file = 'C:/Users/Hp/Desktop/Engineering AI So1/5.2d/extracted_frames/frame_0052.jpg'
2 image = Image.open(image_file)
3 with open(image_file, mode="rb") as image_data:
4     results = prediction_client.detect_image(project_id, model_name, image_data)

In [63]: 1 for obj in results.predictions:
2     if (obj.probability*100) > 30:
3         print("Detected object:", obj.tag_name)
4         print("Bounding box:", obj.bounding_box)
5         print("Confidence:", obj.probability)
6         print()

Detected object: car
Bounding box: {'additional_properties': {}, 'left': 0.24527396, 'top': 0.36195692, 'width': 0.07298212, 'height': 0.1285204}
Confidence: 0.8921732

Detected object: truck
Bounding box: {'additional_properties': {}, 'left': 0.19364765, 'top': 0.2934665, 'width': 0.061764598, 'height': 0.08511239}
Confidence: 0.8501856

Detected object: car
Bounding box: {'additional_properties': {}, 'left': 0.29169586, 'top': 0.2887986, 'width': 0.042362392, 'height': 0.05960539}
Confidence: 0.64105654

Detected object: truck
Bounding box: {'additional_properties': {}, 'left': 0.2937003, 'top': 0.2830731, 'width': 0.04151711, 'height': 0.060719848}
Confidence: 0.5979977

```

Prediction_client sends a request to the Custom Vision prediction service with the provided project ID, model name, and image data. The service processes the image using the specified model and returns the results, which includes information about the detected objects in the image, such as their location and labels. These results are then stored in the results variable for further processing and analysis. In the second block of code the results variable is used and the detected objects information is displayed.

Object detection

To detect the objects I wrote a function that is responsible for detecting objects in a frame using Azure Custom Vision.

The function detect_object, takes a single argument, frame, which represents a single image. The frame is converted to bytes using OpenCV', which encodes the frame as a JPEG image. The resulting encoded image is converted to bytes. This converts the image data into a format that can be sent over the network. The converted image bytes are used as input to call the Custom Vision API to detect objects. The prediction_client.detect_image() function is invoked with the required parameters. The results of the object detection are stored in the results.

The function then proceeds to extract the bounding box information for the detected objects from the results variable. It initializes two empty lists, detections and tag_names, to store the bounding box coordinates and tag names respectively for each detected object. It iterates over each prediction in the results. If the confidence probability of a prediction is above 40%, the bounding box coordinates are calculated using the position and dimensions of the bounding box in the frame. These coordinates are cast to integers. The calculated bounding box coordinates and the tag name are appended to the respective lists.

Finally, the function returns a NumPy array containing the detection and tag_names list. The bounding box coordinates and tag names for each detected object are returned as arrays, allowing for further processing and visualization.

Below is the screen shot of the function detect_object.

```
In [67]: 1 #iteration 2
2 # Function to detect objects in a frame using Azure Custom Vision
3 def detect_objects(frame):
4     # Convert frame to bytes
5     _, img_encoded = cv2.imencode('.jpg', frame)
6     image_bytes = img_encoded.tobytes()
7
8     # Call the Custom Vision API to detect objects
9     results = prediction_client.detect_image(project_id, model_name, image_bytes)
10
11    # Extract bounding box information for detected objects
12    detections = []
13    tag_names = [] # List to store tag names corresponding to each detected object
14    for prediction in results.predictions:
15        if prediction.probability > 0.4: # Only include predictions with a confidence above 40%
16            left = int(prediction.bounding_box.left * frame.shape[1])
17            top = int(prediction.bounding_box.top * frame.shape[0])
18            right = int(left + prediction.bounding_box.width * frame.shape[1])
19            bottom = int(top + prediction.bounding_box.height * frame.shape[0])
20            detections.append([left, top, right, bottom])
21            tag_names.append(prediction.tag_name)
22
23    return np.array(detections), tag_names
```

Object Tracking

In order to track the objects, I decided to make a function that is responsible for tracking objects using the SORT (Simple Online and Realtime Tracking) algorithm. The function, track_objects(), take two arguments, detections and tracking. Detection represents the detected objects in a frame, typically in the form of bounding box coordinates and tracker is an object that implements the SORT algorithm for object tracking.

Inside the function, the tracker_update() method is called with the detections as its argument. This method updates the tracker with the newly detected objects and returns the tracked objects. The tracked objects are stored in the variable tracked_objects, which is then returned.

```
In [68]: 1 # Function to track objects using SORT algorithm
2 def track_objects(detections, tracker):
3     tracked_objects = tracker.update(detections)
4     return tracked_objects|
```

Object Counts

This function helps in visualizing the count of the objects per frame. It gets the object_count and frame which are displayed on to the frame using cv2.putText method.

```
In [69]: 1 # Function to visualize object counts|  
2 def visualize_object_counts(frame, object_count, tag_names):  
3     # Write object count and tag names on the frame  
4     cv2.putText(frame, f"Object Count: {object_count}", (10, 30), cv2.FONT_HERSHEY_SIMPLEX, 0.7, (0, 255, 0), 2)
```

Working of main()

The main function orchestrates the process of detecting, tracking, and visualizing objects within a video stream. It begins by initializing a video capture object, which opens the specified video file for frame-by-frame reading. Concurrently, an instance of the SORT (Simple Online and Realtime Tracking) algorithm is created to facilitate object tracking. Additionally, a VideoWriter object is established to write the processed frames into an output video file. Within the main loop, each frame of the input video is iterated over, incrementing the frame count and breaking the loop if no further frames are available. During each iteration, objects within the frame are detected using the Azure Custom Vision API, followed by tracking these objects using the SORT algorithm. Subsequently, the number of objects detected in each frame is counted, and visualized on the frame. Additionally, bounding boxes and unique IDs are drawn around the tracked objects to enhance visibility, these unique ids help identify and track the object in each frame. Finally, the processed frame is written into the output video file, and the frame is displayed. The program can be terminated by pressing the 'q' key, at which point the video capture and output objects are released, and all OpenCV windows are closed.

Below is the code for calling all the functionalities using main.

```
n [*]: 1 # Main function  
2 def main():  
3  
4     # Initialize video capture  
5     video_capture = cv2.VideoCapture('Highway Traffic.mp4')  
6  
7     # Initialize SORT tracker  
8     tracker = Sort()  
9  
10    # Define the codec and create VideoWriter object  
11    fourcc = cv2.VideoWriter_fourcc(*'XVID')  
12    out = cv2.VideoWriter('output_video.avi', fourcc, 20.0, (int(video_capture.get(3)), int(video_capture.get(4))))  
13  
14    # Loop through frames  
15    frame_count = 0  
16    while video_capture.isOpened():  
17        ret, frame = video_capture.read()  
18        if not ret:  
19            break  
20  
21        # Detect objects in the frame  
22        detections, tag_names = detect_objects(frame)  
23  
24        # Track objects using SORT  
25        tracked_objects = track_objects(detections, tracker)  
26  
27        # Count number of objects per frame  
28        object_count = len(tracked_objects)  
29
```

```

# Visualize object counts and tag names on the frame
visualize_object_counts(frame, object_count, tag_names)

# Draw bounding boxes and IDs for tracked objects
for obj in tracked_objects:
    left, top, right, bottom, track_id = obj
    # Ensure bounding box coordinates are integers
    left = int(left)
    top = int(top)
    right = int(right)
    bottom = int(bottom)
    cv2.rectangle(frame, (left, top), (right, bottom), (0, 255, 0), 2)
    cv2.putText(frame, f"ID: {track_id}", (left, top - 10), cv2.FONT_HERSHEY_SIMPLEX, 0.5, (0, 255, 0), 2)

# Write the frame into the output video
out.write(frame)

# Display the frame
cv2.imshow('Video', frame)

```

```

50
51      # Break the loop if 'q' key is pressed
52      if cv2.waitKey(1) & 0xFF == ord('q'):
53          break
54
55      # Release the video capture object
56      video_capture.release()
57      out.release()
58      cv2.destroyAllWindows()
59
60
61  if __name__ == "__main__":
62      main()
63
64
65
66
67

```

Below is the screenshot of the detected objects, their tracking ID and the object count of one of the frames.

